

Assignment 2

Haiyu Wang

1. Configuration File:

Assignment2.py

```
# Assignment 2, Haiyu Wang
```

```
# Implement Cache and test the performance
```

```
import m5
```

```
from m5.objects import *
```

```
import os
```

```
gem5_path = os.environ["GEM5"]
```

```
import optparse
```

```
parser = optparse.OptionParser()
```

```
parser.add_option("--prog", type="str", default=None);
```

```
parser.add_option("--clock_freq", type="str", default=None);
```

```
parser.add_option("--l1i_size", type="str", default=None);
```

```
parser.add_option("--l1i_assoc", type="int", default=None);
```

```
parser.add_option("--l1d_size", type="str", default=None);
```

```
parser.add_option("--l1d_assoc", type="int", default=None);
```

```
(options, args) = parser.parse_args()
```

```
program = options.prog
```

```
system = System()
```

```
system.clk_domain = SrcClockDomain()
```

```
system.clk_domain.voltage_domain = VoltageDomain()
```

```
isa = m5.defines.buildEnv['TARGET_ISA']
```

```

if options.clock_freq:
    system.clk_domain.clock = options.clock_freq
else:
    system.clk_domain.clock = '1.2GHz'

from Caches import *

system.mem_mode = 'timing'
system.mem_ranges = [AddrRange('512MB')]

system.cpu = TimingSimpleCPU()

system.membus = SystemXBar()

system.cpu.icache = L1ICache(options)
system.cpu.dcache = L1DCache(options)

system.cpu.icache.connectCPU(system.cpu)
system.cpu.dcache.connectCPU(system.cpu)

system.l2bus = L2XBar()
system.cpu.icache.connectBus(system.l2bus)
system.cpu.dcache.connectBus(system.l2bus)

system.l2cache = L2Cache(options)
system.l2cache.connectCPUSideBus(system.l2bus)
system.l2cache.connectMemSideBus(system.membus)

system.cpu.createInterruptController()
if isa == 'x86':
    system.cpu.interrupts[0].pio = system.membus.master
    system.cpu.interrupts[0].int_master = system.membus.slave
    system.cpu.interrupts[0].int_slave = system.membus.master

```

```

system.system_port = system.membus.slave

system.mem_ctrl = DDR3_1600_8x8()
system.mem_ctrl.range = system.mem_ranges[0]
system.mem_ctrl.port = system.membus.master

process = Process()
apps_path = "/project/linuxlab/gem5/test_progs"
if program == "daxpy" and isa == "x86":
    process.cmd = [apps_path + '/daxpy/daxpy_x86']
elif program == "daxpy" and isa == "arm":
    process.cmd = [apps_path + '/daxpy/daxpy_arm']
elif program == "queens" and isa == "x86":
    process.cmd = [apps_path + '/queens/queens_x86']
    process.cmd += ["10 -c"]
elif program == "queens" and isa == "arm":
    process.cmd = [apps_path + '/queens/queens_arm']
    process.cmd += ["10 -c"]

system.cpu.workload = process
system.cpu.createThreads()

root = Root(full_system = False, system = system)
m5.instantiate()
print ("Beginning simulation!")

exit_event = m5.simulate()
print('Exiting @ tick %i because %s' % (m5.curTick(),
exit_event.getCause()))

```

Caches.py

```
from m5.defines import buildEnv
from m5.objects import *

class L1Cache(Cache):
    assoc = 2
    tag_latency = 2
    data_latency = 2
    response_latency = 2
    mshrs = 4
    tgts_per_mshr = 20

    def __init__(self, options=None):
        super(L1Cache, self).__init__()
        pass

    def connectBus(self, bus):
        self.mem_side = bus.slave

    def connectCPU(self, cpu):
        raise NotImplementedError

class L1ICache(L1Cache):
    is_read_only = True
    writeback_clean = True
    size = '16kB'

    def __init__(self, opts=None):
        super(L1ICache, self).__init__(opts)
        if opts.l1i_size:
            self.size = opts.l1i_size
        if opts.l1i_assoc:
            self.assoc = opts.l1i_assoc

    def connectCPU(self, cpu):
```

```
self.cpu_side = cpu.icache_port
```

```
class L1DCache(L1Cache):  
    is_read_only = False  
    writeback_clean = False  
    size = '64kB'  
  
    def __init__(self,opts=None):  
        super(L1DCache,self).__init__(opts)  
        if opts.l1d_size:  
            self.size = opts.l1d_size  
        if opts.l1d_assoc:  
            self.assoc = opts.l1d_assoc  
  
    def connectCPU(self, cpu):  
        self.cpu_side = cpu.dcache_port
```

```
class L2Cache(Cache):  
    is_read_only = False  
    writeback_clean = False  
    size = '256kB'  
    assoc = 8  
    tag_latency = 20  
    data_latency = 20  
    response_latency = 80  
    mshrs = 20  
    tgts_per_mshr = 12  
  
    def __init__(self,options=None):  
        super(L2Cache,self).__init__()  
        pass  
  
    def connectCPUSideBus(self,bus):
```

```

self.cpu_side = bus.master

def connectMemSideBus(self,bus):

    self.mem_side = bus.slave

```

2. Console Output:

```

[haiyu.wang@shell hw2]$ $GEM5/build/ARM/gem5.opt --outdir="daxpy_arm" assignment2.py --prog="daxpy"
gem5 Simulator System.  http://gem5.org
gem5 is copyrighted software; use the --copyright option for details.

gem5 compiled Sep  9 2019 15:01:37
gem5 started Oct 29 2019 15:18:36
gem5 executing on shell.cec.wustl.edu, pid 2292
command line: /project/linuxlab/gem5/gem5_dev/build/ARM/gem5.opt --outdir=daxpy_arm assignment2.py --prog=daxpy

Global frequency set at 1000000000 ticks per second
warn: DRAM device capacity (8192 Mbytes) does not match the address range assigned (512 Mbytes)
0: system.remote_gdb: listening for remote gdb on port 7000
Beginning simulation!
info: Entering event queue @ 0.  Starting simulation...
info: Increasing stack size by one page.
warn: readlink() called on '/proc/self/exe' may yield unexpected results in various settings.
    Returning '/project/linuxlab/gem5/test_progs/daxpy/daxpy_arm'
7425.000000
Exiting @ tick 111228824 because exiting with last active thread context
[haiyu.wang@shell hw2]$

```


3. Graphs and Tables

a) L1D Overall Hits

i. Table

L1D Overall Hits		
	1.0 GHz, 1 Association	0.8 GHz, 2 Association
8kB	3173	3207
16kB	3221	3230
32kB	3240	3243
64kB	3243	3250

ii. Graph

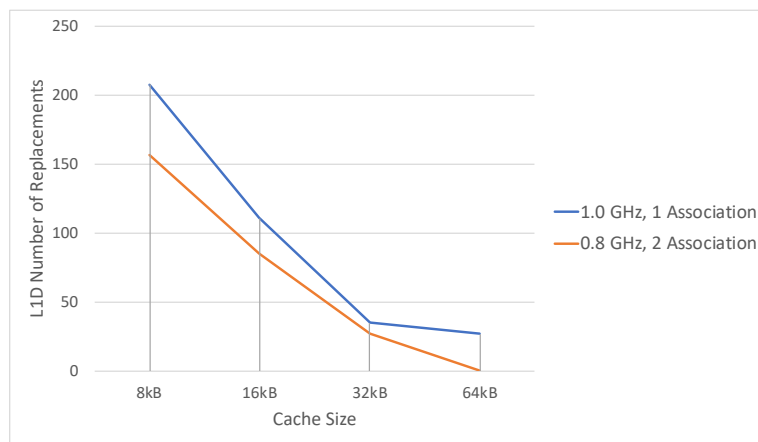


b) L1D Number of Replacements

i. Table

L1D Number of Replacements		
	1.0 GHz, 1 Association	0.8 GHz, 2 Association
8kB	208	157
16kB	111	85
32kB	35	27
64kB	27	0

ii. Graph

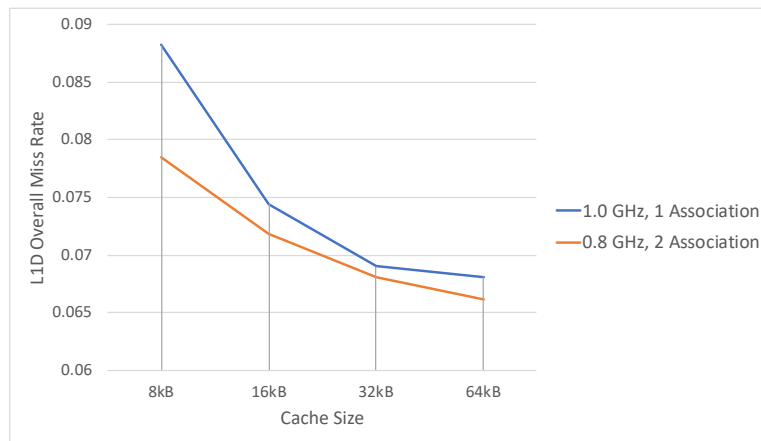


c) L1D Overall Miss Rate

i. Table

L1D Overall Miss Rate		
	1.0 GHz, 1 Association	0.8 GHz, 2 Association
8kB	0.088218	0.078448
16kB	0.074425	0.071839
32kB	0.068996	0.068103
64kB	0.068103	0.066092

ii. Graph




d) L2 Overall Miss Rate

i. Table

L2 Overall Miss Rate		
	1.0 GHz, 1 Association	0.8 GHz, 2 Association
8kB	0.850136	0.891429
16kB	0.909621	0.921713
32kB	0.935532	0.939759
64kB	0.939759	0.949772

ii. Graph



4. Answers

CPI = number of Cycles / number of instructions.

Case 1:

1.0 GHz, 1 Association, 8KB: CPI = 13.84

1.0 GHz, 1 Association, 16KB: CPI = 13.72

In this case, the cache size doubles, and the instructions committed remain, but the number of CPU cycles in the bigger cache drops (because it can cache

more data, so the number of hits increase which leads to decreasing of number of CPU cycles), so CPI drops.

Case 2:

1.0 GHz, 1 Association, 8KB: CPI = 13.84

0.8 GHz, 2 Association, 8KB: CPI = 13.08

In this case, although the case sizes are the same, the latter one has 2 associations which means there are more hits when running the programs.

Thus, the total cycles will decrease, which leads to the decreasing of CPI