

CSE 560 Computer Systems Architecture

Static Scheduling

1

Multiple Issue Redux

Which of the following statements is false?

- A. Multiple issue is needed to expose insn level parallelism (ILP) beyond pipelining
- B. Multiple issue improves performance
- C. Multiple issue improves utilization
- D. Multiple issue doesn't make sense past 6-wide
- E. Multiple issue can be a hardware or a software technique



2

Multiple Issue Redux


Which of the following statements is false?

- A. Multiple issue is needed to expose insn level parallelism (ILP) beyond pipelining
- B. Multiple issue improves performance
- C. **Multiple issue improves utilization**
- D. Multiple issue doesn't make sense past 6-wide
- E. Multiple issue can be a hardware or a software technique

3

3

This Unit: Static Scheduling



- Code scheduling to
 - Reduce pipeline stalls
 - Increase ILP

Two approaches to scheduling

- **This Unit:**
 - **Static scheduling by the compiler**
- Coming Soon:
 - Dynamic scheduling by the hardware

4

4

Code Scheduling

- Scheduling: act of finding independent instructions
 - **Static:** at compile time by the compiler (software)
 - **Dynamic:** at runtime by the processor (hardware)
- Why schedule code?
 - **Scalar pipelines:** fill load-to-use delays to improve CPI
 - **Superscalar:** place independent instructions together
 - As above, load-to-use delay slots
 - Allow multiple-issue decode logic to let them execute at the same time

5

5

Scheduling Requirements

- **Independent insns**
 - no ILP → game over
- **Large Scheduling Scope**
 - Scope = code region we are scheduling
 - The bigger the better (more independent insns to play with)
 - Once scope is defined, schedule is pretty obvious
 - Trick is creating a large scope (schedule across branches?)
- **Enough registers**
 - To hold additional "live" values
- **Alias analysis**
 - Whether load/store reference same memory locations
 - Can they be legally rearranged?

6

6

Scheduling Techniques

- **Stall Removal**
 - Separate load-use pairs
- **Scope enlarging**
 - For Loops: **loop unrolling**
 - For Non-loops:
 - **Superblocks**
 - **Predication**
- **Exploit Data-Level Parallelism**
 - Vectors

7

New Metric: Utilization

- Utilization:** actual performance / peak performance
- Important metric for performance/cost
 - Why pay for hardware you rarely use?
 - Adding hardware usually \uparrow performance, \uparrow utilization
 - New hardware cannot always be exploited
 - Diminishing marginal returns
 - Compiler can help make better use of existing hardware
 - Important for superscalar

8

Running Code Example: SAXPY

- **SAXPY** (Single-precision A X Plus Y)
 - Linear algebra routine (for solving systems of equations)
 - Part of early **Livermore Loops** benchmark suite
 - floating point uses "F" registers and "F" instructions

```
for (i=0; i<N; i++)
    Z[i] = (A*X[i]) + Y[i];
```

```
0: ldf X(r1) → f1      // loop          LOAD1
1: mulf f0, f1 → f2    // A in f0      USE1
2: ldf Y(r1) → f3      // X, Y, Z constants  LOAD2
3: addf f2, f3 → f4    USE2
4: stf f4 → Z(r1)
5: addi r1, 4 → r1     // i in r1
6: blt r1, r2, 0       // N*4 in r2
```

9

SAXPY Performance and Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1) → f1	F	D	X	M	W															
mulf f0, f1 → f2		F	D	d*	E*	E*	E*	E*	W											
ldf Y(r1) → f3			F	p*	D	X	M	W												
addf f2, f3 → f4				F	d*	d*	D	E+	E+	W										
stf f4 → Z(r1)					F	p*	p*	p*	D	X	M	W								
addi r1, 4 → r1									F	D	X	M	W							
blt r1, r2, 0										F	D	X	M	W						
ldf X(r1) → f1											F	D	X	M	W					

Scalar pipeline

- Full bypassing, 5-cycle E*, 2-cycle E+, predict branches taken
- Single iteration (7 insns) latency: **16-5 = 11 cycles**
- **Performance:** 7 insns / 11 cycles = 0.64 IPC
- **Utilization:** actual/peak IPC = 0.64 / 1 = 64%

10

SAXPY Performance and Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1) → f1	F	D	X	M	W															
mulf f0, f1 → f2		F	D	d*	E*	E*	E*	E*	W											
ldf Y(r1) → f3			F	D	X	M	W													
addf f2, f3 → f4			F	p*	p*	d*	d*	D	E+	E+	W									
stf f4 → Z(r1)				F	p*	p*	p*	p*	p*	D	X	M	W							
addi r1, 4 → r1					F	p*	p*	p*	p*	D	X	M	W							
blt r1, r2, 0						F	p*	p*	p*	p*	D	X	M	W						
ldf X(r1) → f1											F	D	X	M	W					

2-way superscalar pipeline

- Any two insns per cycle + split integer and FP pipelines
- + **Performance:** 7 insns / 10 cycles = 0.70 IPC
- **Utilization:** actual/peak IPC = 0.70 / 2 = 35%
- More hazards → more stalls
- Each stall is more expensive

11

Eliminate Load-Use Pairs?

```
for (i=0; i<N; i++)
    Z[i] = (A*X[i]) + Y[i];
```

```
0: ldf X(r1) → f1      LOAD1
1: mulf f0, f1 → f2    USE1
2: ldf Y(r1) → f3      LOAD2
3: addf f2, f3 → f4    USE2
4: stf f4 → Z(r1)
5: addi r1, 4 → r1
6: blt r1, r2, 0

0: ldf X(r1) → f1      LOAD1
2: ldf Y(r1) → f3      LOAD2
1: mulf f0, f1 → f2    USE1
3: addf f2, f3 → f4    USE2
4: stf f4 → Z(r1)
5: addi r1, 4 → r1
6: blt r1, r2, 0
```

Problem solved?

12

Loop Unrolling SAXPY

- **Goal:** separate dependent insns from one another
- SAXPY problem: not enough flexibility within one iteration
 - Longest chain of insns is 9 cycles
 - Load (1)
 - Forward to multiply (5)
 - Forward to add (2)
 - Forward to store (1)
 - Can't hide a 9-cycle chain using only 7 insns
 - But how about two 9-cycle chains using 14 insns?
- **Loop unrolling:** schedule 2+ iterations together
 - Fuse iterations
 - Schedule to reduce stalls
 - Schedule introduces ordering problems → rename registers

13

Unrolling SAXPY I: Fuse Iterations

- Combine two (in general K) iterations of loop
 - Fuse loop control: induction variable ($i=r1$) increment + branch
 - Adjust (implicit) induction uses: constants → constants + 4

```
ldf X(r1),f1      ldf X(r1),f1
mul f0,f1,f2      mul f0,f1,f2
ldf Y(r1),f3      ldf Y(r1),f3
addf f2,f3,f4      addf f2,f3,f4
stf f4,Z(r1)      stf f4,Z(r1)
addi r1,4,r1      addi r1,4,r1
blt r1,r2,0      blt r1,r2,0
-- increment i    -- increment i
-- jump back      -- jump back

ldf X+4(r1),f1
mul f0,f1,f2
ldf Y+4(r1),f3
addf f2,f3,f4
stf f4,Z+4(r1)
addi r1,8,r1
blt r1,r2,0
```

14

Unrolling SAXPY II: Pipeline Schedule

- Pipeline schedule to reduce stalls
 - Have already seen this: pipeline scheduling

```
ldf X(r1),f1      ldf X(r1),f1
mul f0,f1,f2      ldf X+4(r1),f1
ldf Y(r1),f3      mul f0,f1,f2
addf f2,f3,f4      mul f0,f1,f2
stf f4,Z(r1)      ldf Y(r1),f3
ldf X+4(r1),f1      ldf Y+4(r1),f3
mul f0,f1,f2      addf f2,f3,f4
ldf Y+4(r1),f3      addf f2,f3,f4
addf f2,f3,f4      stf f4,Z(r1)
stf f4,Z+4(r1)      stf f4,Z+4(r1)
addi r1,8,r1      addi r1,8,r1
blt r1,r2,0      blt r1,r2,0
```

15

Unrolling SAXPY III: "Rename" Registers

- Pipeline scheduling causes reordering violations
 - Rename registers to correct

```
ldf X(r1),f1      ldf X(r1),f1
ldf X+4(r1),f1      ldf X+4(r1),f5
mul f0,f1,f2      mul f0,f1,f2
mul f0,f1,f2      mul f0,f5,f6
ldf Y(r1),f3      ldf Y(r1),f3
ldf Y+4(r1),f3      ldf Y+4(r1),f7
addf f2,f3,f4      addf f2,f3,f4
addf f2,f3,f4      addf f6,f7,f8
stf f4,Z(r1)      stf f4,Z(r1)
stf f4,Z+4(r1)      stf f8,Z+4(r1)
addi r1,8,r1      addi r1,8,r1
blt r1,r2,0      blt r1,r2,0
```

Are we sure we can move these loads above these stores?

Alias analysis must be conservative.

16

Unrolled SAXPY Performance/Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1),f1	F	D	X	M	W															
ldf X+4(r1),f5		F	D	X	M	W														
mul f0,f1,f2			F	D	E*	E*	E*	E*	W											
mul f0,f5,f6				F	D	E*	E*	E*	E*	W										
ldf Y(r1),f3					F	D	X	M	W											
ldf Y+4(r1),f7						F	D	X	M	W										
addf f2,f3,f4							F	D	X	M	W									
addf f6,f7,f8								F	D	X	M	W								
stf f4,Z(r1)									F	D	X	M	W							
stf f8,Z+4(r1)										F	D	X	M	W						
addi r1,8,r1											F	D	X	M	W					
blt r1,r2,0												F	D	X	M	W				
ldf X(r1),f1													F	D	X	M	W			

+ Performance: 12 insn / 13 cycles = 0.92 IPC

+ Utilization: actual/peak IPC = 0.92 / 1 = 92%

+ **Speedup:** (2 * 11 cycles) / 13 cycles = 1.69

? But improvement in IPC is only 0.92/0.64 = 1.43, what gives? 17

Loop Unrolling Shortcomings

- Static code growth → more I\$ misses (limits unrolling)
- Needs more registers to hold values (ISA limits this)
- Doesn't handle: non-loops, inter-iteration dependences

```
for (i=0; i<N; i++)
  X[i]=A*X[i-1];
```

```
ldf X-4(r1),f1      ldf X-4(r1),f1
mul f0,f1,f2      mul f0,f1,f2
stf f2,X(r1)      stf f2,X(r1)
addi r1,4,r1      mul f0,f2,f3
blt r1,r2,0      stf f3,X+4(r1)
ldf X-4(r1),f1      addi r1,4,r1
mul f0,f1,f2      blt r1,r2,0
stf f2,X(r1)
```

- Two mul's are not parallel
- Other (more advanced) techniques help

18

13

14

15

16

17

18

Summary: Static Scheduling Limitations

- Limited number of registers (set by ISA)
- Scheduling scope
 - Example: hard to move memory insns past branches
- Inexact memory aliasing information
 - Often prevents reordering of loads above stores
- Caches misses (or any runtime event) confound scheduling
 - How can the compiler know which loads will miss/hit?
 - Can impact the compiler's scheduling decisions

19

19

Scheduling Techniques

- Stall Removal**
 - Separate load-use pairs
- Scope enlarging**
 - For Loops: **loop unrolling**
 - For Non-loops:**
 - Superblocks** (biased branches)
 - Predication** (non-biased branches)
- Exploit Data-Level Parallelism**
 - Vectors

20

20

Superblocks


Source code

```
A = Y[i];
if (A == 0)
    A = W[i];
else
    Y[i] = 0;
Z[i] = A*X[i];
```

Machine code

```
0: ldf Y(r1), f2
1: fbne f2, 4
2: ldf W(r1), f2
3: jump 5
4: stf f0, Y(r1)
5: ldf X(r1), f4
6: mul f4, f2, f6
7: stf f6, Z(r1)
```

4 basic blocks: A, B, C, D



- Use when branch is highly biased
- Fuse blocks of most frequent path: ACD
- Schedule
- Create **repair code** in case path = ABD

21

21

Superblock and Repair Code

Superblock

```
0: ldf Y(r1), f2
1: fbeq f2, 2
4: stf f0, Y(r1)
5: ldf X(r1), f4
6: mul f4, f2, f6
7: stf f6, Z(r1)
```

Repair code

```
2: ldf W(r1), f2
5': ldf X(r1), f4
6': mul f4, f2, f6
7': stf f6, Z(r1)
```

- What did we do?
 - Change sense (test) of branch 1
 - Original taken target now fall-thru
 - Created repair block
 - May need to duplicate some code (here basic-block D)
 - Haven't actually scheduled superblock yet

22

22

Superblocks Scheduling I

Superblock

```
0: ldf Y(r1), f2
1: fbeq f2, 2
5: ldf X(r1), f4
6: mul f4, f2, f6
4: stf f0, Y(r1)
7: stf f6, Z(r1)
```

Repair code

```
2: ldf W(r1), f2
5': ldf X(r1), f4
6': mul f4, f2, f6
7': stf f6, Z(r1)
```

- First scheduling move: move insns 5 and 6 above insn 4
 - Hmmm: moved load (5) above store (4)
 - We can tell this is OK, but can the compiler
 - If yes, fine
 - Otherwise, need to do something

23

23

Predication

- Conventional control
 - Conditionally executed insns also conditionally fetched
- Predication**
 - Conditionally executed insns unconditionally fetched
 - Full predication** (ARM, IA-64)
 - Tag every insn with predicate, costs extra bits
 - Conditional moves** (Alpha, IA-32)
 - Construct appearance of full predication from one primitive


```
cmovneq r1, r2, r3 // if (r1==0) r3 ← r2;
```

 - May require some code duplication to achieve desired effect
 - Only good way of adding predication to an existing ISA
 - If-conversion:** replacing control with predication
 - + Good if branch is unpredictable (save mis-prediction)
 - But more instructions fetched and "executed"

24

24

Avoiding Branches via ISA: Predication

- Conventional control
 - Conditionally executed insns also conditionally fetched
- | | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|----------------------|---|---|---|---|----|----|----|---------------------|---|
| beq r3, targ | F | D | X | M | W | | | | |
| sub r6, 1, r5 | | | F | D | -- | -- | -- | flushed: wrong path | |
| targ: add r4, r5, r4 | | | | F | -- | -- | -- | flushed: why? | |
| targ: add r4, r5, r4 | | | | | F | D | X | M | W |
- If **beq** mis-predicts, both **sub** and **add** must be flushed
 - Waste: **add** is independent of mis-prediction
 - Predication**: not prediction, predication
 - ISA support for conditionally-executed unconditionally-fetched insns
 - If **beq** mis-predicts, annul **sub** in place, preserve **add**
 - Example is if-then, but if-then-else can be predicated too
 - How is this done? How does **add** get correct value for **r5**

25

25

Full Predication

Full predication

- Every insn can be annulled, annulment controlled by...
- Predicate registers: additional register in each insn (e.g., IA64)

	1	2	3	4	5	6	7	8	9
setp.eq r3, p3	F	D	X	M	W				
sub.p r6, 1, r5, p3			F	D	X	--	--	annulled	
targ: add r4, r5, r4				F	D	X	M	W	

- Predicate codes: condition bits in each insn (e.g., ARM)

	1	2	3	4	5	6	7	8	9
setcc r3	F	D	X	M	W				
sub.nz r6, 1, r5			F	D	X	--	--	annulled	
targ: add r4, r5, r4				F	D	X	M	W	

- Only ALU insn shown (**sub**), but this applies to all insns, even stores
- Branches replaced with "set-predicate" insns

26

26

Conditional Register Moves (CMOVs)

Conditional (register) moves


- Construct appearance of full predication from one primitive
 - cmovneq r1, r2, r3** // if (r1==0) r3 ← r2;
- May require some code duplication to achieve desired effect
- Painful, potentially impossible for some insn sequences
- Requires more registers
- Only good way of retro-fitting predication onto ISA (e.g., IA32, Alpha)

	1	2	3	4	5	6	7	8	9
sub r6, 1, r9			D	X	M	W			
cmovne r3, r9, r5			F	D	X	M	W		
targ: add r4, r5, r4				F	D	X	M	W	

27

27

Non-Biased Branches: Use Predication



Using Predication

0: ldf Y(r1), f2
1: fspne f2, p1
2: ldf.p p1, W(r1), f2
4: stf.np p1, f0, Y(r1)
5: ldf X(r1), f4
6: mul f4, f2, f6
7: stf f6, 2(r1)

28

28

ISA Support for Predication

0: ldf Y(r1), f2
1: fspne f2, p1
2: ldf.p p1, W(r1), f2
4: stf.np p1, f0, Y(r1)
5: ldf X(r1), f4
6: mul f4, f2, f6
7: stf f6, 2(r1)

- IA-64: change branch 1 to **set-predicate insn fspne**
- Change insns 2 and 4 to **predicated insns**
 - ldf.p** performs **ldf** if predicate **p1** is true
 - stf.np** performs **stf** if predicate **p1** is false

29

29

Predication Performance

- Cost/benefit analysis
 - Benefit**: predication avoids branches
 - Thus avoiding mis-predictions
 - Also reduces pressure on predictor table (few branches to track)
 - Cost**: extra (annulled) instructions
- Since branch predictors are highly accurate...
 - Might not help:
 - 5-stage pipeline, two instruction on each path of if-then-else
 - No performance gain, likely slower if branch predictable
 - Or even hurt!
 - But can help:
 - Deeper pipelines, hard-to-predict branches, and few added insns
- Predication is useful, but not a panacea

30

30

Aside: Profiling

How do we know whether a branch is biased or not?

- Profile:** statistical information about program tendencies
 - Collect from previous program runs (different inputs)
- ± Works OK depending on information
 - Memory latencies (cache misses)
 - + Which loads miss frequently independent of inputs?
 - Depends on cache configuration
 - Memory dependences
 - + Which loads & stores communicate with each other?
 - + Stable across inputs
 - Branch outcomes
 - + Which branches are usually taken/not-taken?
 - Not so stable across inputs
- Popular research topic

31

31

Scheduling Techniques

- Stall Removal**
 - Separate load-use pairs
- Scope enlarging**
 - For Loops: **loop unrolling**
 - For Non-loops:**
 - Superblocks** (biased branches)
 - Predication** (non-biased branches)
- Exploit Data-Level Parallelism**
 - Vectors

32

32

Data-Level Parallelism

Data-level parallelism (DLP)

- Single operation repeated on multiple data elements
 - SIMD (Single-Instruction, Multiple-Data)
- Less general than ILP: parallel insns are all same operation
- Exploit with **vectors**

Old idea: Cray-1 supercomputer from late 1970s

- Eight 64-entry x 64-bit floating point "Vector registers"
 - 4096 bits (0.5KB) in each register! 4KB vector register file
- Special vector instructions to perform vector operations
 - Load vector, store vector (wide memory operation)
 - Vector+Vector addition, subtraction, multiply, etc.
 - Vector+Constant addition, subtraction, multiply, etc.
 - In Cray-1, each instruction specifies 64 operations!

33

33

Example Vector ISA Extensions

Extend ISA with floating point (FP) vector storage ...

- Vector register:** fixed-size array of 32- or 64- bit FP elements
- Vector length:** For example: 4, 8, 16, 64, ...
- ... and example operations for vector length of 4, 8-bit elements
 - Load vector: `ldf.v X(r1), v1` =


```
ldf X+0(r1), v1[0]
ldf X+1(r1), v1[1]
ldf X+2(r1), v1[2]
ldf X+3(r1), v1[3]
```
 - Add two vectors: `addf.vv v1, v2, v3`

```
addf v1[i], v2[i], v3[i] (where i is 0,1,2,3)
```
 - Add vector to scalar: `addf.vs v1, f2, v3`

```
addf v1[i], f2, v3[i] (where i is 0,1,2,3)
```

34

34

Example Use of Vectors – 4-wide

7x1024 insns	<pre>ldf X(r1), f1 mulf f0, f1, f2 ldf Y(r1), f3 addf f2, f3, f4 stf f4, Z(r1) addi r1, 4, r1 blti r1, 4096, 0</pre>		7x256 insns (4x fewer insns)	<pre>ldf.v X(r1), v1 mulf.vs v1, f0, v2 ldf.v Y(r1), v3 addf.vv v2, v3, v4 stf.v v4, Z(r1) addi r1, 16, r1 blti r1, 4096, 0</pre>
--------------	--	--	------------------------------	---

Operations

- Load vector: `ldf.v X(r1), v1`
- Multiply vector to scalar: `mulf.vs v1, f2, v3`
- Add two vectors: `addf.vv v1, v2, v3`
- Store vector: `stf.v v1, X(r1)`

Performance?

- If CPI = 1, 4x speedup
- CPI not always 1
 - Execution width (implementation) ≠ vector width (ISA)

35

35

Why Vectorization is Awesome

Have your cake and eat it, too

All the benefits of a wider machine, without superscalar costs

- Single instruction fetch
- Wide reads & writes (without multiple \$ or regfile ports)
- Wider data to bypass ≠ N² bypass

Execution width (implementation) vs vector width (ISA)

- Example: Pentium 4 and Core 1 execute vector ops at half width
- Core 2 executes them at full width
- Intel's Sandy Bridge brings 256-bit vectors to x86
- Intel's Larrabee graphics chip brings 512-bit vectors to x86

Vector + superscalar? Sure!

- Multiple n-wide vector instructions per cycle

36

36

Scheduling: Compiler or Hardware

Compiler

- + Large scheduling scope (full program)
- + Simple hardware → fast clock, short pipeline, and low power
- Low branch prediction accuracy (profiling?)
- Little information on memory dependences (profiling?)
- Can't dynamically respond to cache misses (or anything really)
- Hard to speculate, recover from mis-speculation (h/w support?)

Hardware

- Finite buffering resources fundamentally limit scheduling scope
- Scheduling machinery adds pipeline stages and consumes power
- + High branch prediction accuracy
- + Dynamic information about memory dependences
- + Can respond to cache misses
- + Easy to speculate and recover from mis-speculation

37