


CSE 560 Computer Systems Architecture

Dynamic Scheduling

Slides originally developed by Drew Hilton (IBM)
and Milo Martin (University of Pennsylvania)

1

This Unit: Dynamic Scheduling



- Code scheduling
 - To reduce pipeline stalls
 - To increase ILP (insn level parallelism)

Two approaches to scheduling

- Last Unit:
 - Static scheduling by the compiler
- This Unit:
 - Dynamic scheduling by the hardware

2

Scheduling: Compiler or Hardware

Compiler

- + Potentially large scheduling scope (full program)
- + Simple hardware → fast clock, short pipeline, and low power
- Low branch prediction accuracy (profiling?)
- Little information on memory dependences (profiling?)
- Can't dynamically respond to cache misses
- Pain to speculate and recover from mis-speculation (h/w support?)

Hardware

- + High branch prediction accuracy
- + Dynamic information about memory dependences
- + Can respond to cache misses
- + Easy to speculate and recover from mis-speculation
- Finite buffering resources fundamentally limit scheduling scope
- Scheduling machinery adds pipeline stages and consumes power

3

3

Can Hardware Overcome These Limits?

• Dynamically-scheduled processors

- Also called "out-of-order" processors
- Hardware re-schedules insns...
- ...within a sliding window of VonNeumann insns
- As with pipelining and superscalar, ISA unchanged
 - Same hardware/software interface, appearance of in-order
- Increases scheduling scope
 - Does loop unrolling transparently
 - Uses branch prediction to "unroll" branches
- Examples: Pentium Pro/II/III (3-wide), Core 2 (4-wide), Alpha 21264 (4-wide), MIPS R10000 (4-wide), Power5 (5-wide)
- Basic overview of approach

4

4

The Problem With In-Order Pipelines

		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
addf	#0, #1 → #2	F	D	E+	E+	W											
mulf	#2, #3 → #2	F	D*	D*	D	E*	E*	E*	E*	W							
subf	#0, #1 → #4	F	p*	p*	D	E+	E+	E+	W								

- What's happening in cycle 4?
 - **mulf** stalls due to **data dependence**
 - OK, this is a fundamental problem
 - **subf** stalls due to **pipeline hazard**
 - Why? **subf** can't proceed into D because **mulf** is there
 - That is the only reason, and it isn't a fundamental one
 - Maintaining in-order writes to reg. file (both write #2)
- Why can't **subf** go into D in cycle 4 and E+ in cycle 5?

5

5

A Word About Data Hazards

- Real insn sequences pass values via registers/memory
- Three kinds of **data dependences** (where's the fourth?)

	Read-after-write (RAW) True-dependence	Write-after-read (WAR) Anti-dependence	Write-after-write (WAW) Output-dependence
R	add r2, r3 → r1	add r2 , r3 → r1	add r2, r3 → r1
E	sub r1 , r4 → r2	sub r5, r4 → r2	sub r1, r4 → r2
G	or r6, r3 → r1	or r6, r3 → r1	or r6, r3 → r1
M	st r1 → [r2]	ld [r1] → r2	st r1 → [r2]
E	ld [r2] → r4	st r3 → [r1]	st r3 → [r2]
M			

- Only one dependence between any two insns (RAW has priority)
- Focus on **RAW dependences**
- WAR and WAW: less common, just bad naming luck
 - Eliminated by using new register names, (can't rename memory!)

6

6

Find the RAW, WAR, and WAW dependencies

```
add r1 ← r2, r3
sub r4 ← r1, r5
and r2 ← r4, r7
xor r10 ← r2, r11
or r12 ← r10, r13
mult r1 ← r10, r13
```



7

Find the RAW, WAR, and WAW dependencies

```
add r1 ← r2, r3
sub r4 ← r1, r5
and r2 ← r4, r7
xor r10 ← r2, r11
or r12 ← r10, r13
mult r1 ← r10, r13
```

RAW dependencies:

- **r1** from **add** to **sub**
- **r2** from **and** to **xor**
- **r10** from **xor** to **or**
- **r10** from **xor** to **mult**

WAR dependencies:

- **r2** from **add** to **and**
- **r1** from **sub** to **mult**

WAW dependencies:

- **r1** from **add** to **mult**

8

Code Example

- Raw insns:

True Dependencies	False Dependencies
add r2, r3 → r1	add r2, r3 → r1
sub r2, r1 → r3	sub r2, r1 → r3
mul r2, r3 → r3	mul r2, r3 → r3
div r1 , 4 → r1	div r1 , 4 → r1
- "True" (real) & "False" (artificial) dependencies
- Divide insn independent of subtract and multiply insns
 - Can execute in parallel with subtract
- Many registers re-used
 - Just as in static scheduling, the register names get in the way
 - How does the hardware get around this?
- Approach: (step #1) rename registers, (step #2) schedule

9

Step #1: Register Renaming


- To eliminate register conflicts/hazards
- **Architected** vs. **Physical** registers – level of indirection
 - Names: **r1, r2, r3**
 - Locations: **p1, p2, p3, p4, p5, p6, p7**
 - Original mapping: **r1**→**p1, r2**→**p2, r3**→**p3, p4**–**p7** are available

MapTable	FreeList	Original insns	Renamed insns
r1 r2 r3	p4, p5, p6, p7	add r2 , r3 , r1	add p2 , p3 , p4
p1 p2 p3	p5, p6, p7	sub r2 , r1 , r3	sub p2 , p4 , p5
p4 p2 p5	p6, p7	mul r2 , r3 , r3	mul p2 , p5 , p6
p4 p2 p6	p7	div r1 , 4, r1	div p4 , 4, p7
p7 p2 p6			

- Renaming: conceptually write each register once
 - + Removes **false** dependencies
 - + Leaves **true** dependencies intact!
- When to reuse a physical register? After overwriting insn done

10


Step #2: Dynamic Scheduling



- Instructions fetch/decoded/rename into *Instruction Buffer*
 - AKA "instruction window" or "instruction scheduler"
- Instructions (conceptually) check ready bits every cycle
 - Execute when ready

11

Out-of-order Pipeline



12

11

12

REGISTER RENAMING

13

Register Renaming Algorithm

- Data structures:
 - maptable[architectural_reg] → physical_reg
 - Free list: get/put free register
- Algorithm: at decode for each instruction:


```
insn.phys_input1 = maptable[insn.arch_input1]
insn.phys_input2 = maptable[insn.arch_input2]
insn.phys_to_free = maptable[arch_output]
new_reg = get_free_phys_reg()
insn.phys_output = new_reg
maptable[arch_output] = new_reg
```
- At "commit"
 - Once all older instructions have committed, free register
put_free_phys_reg(insn.phys_to_free)

14

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

15

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 →
```

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

16

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
```

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

17

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
```

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

18

17

18

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 →
```

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

19

19

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
sub r5, r2 → r3
addi r3, 1 → r1
```

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

20

20

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
```

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

21

21

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
sub p5, p2 →
addi r3, 1 → r1
```

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

22

22

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
sub p5, p2 → p8
```

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

23

23

Renaming example

Original insns

```
xor r1, r2 → r3
add r3, r4 → r4
sub r5, r2 → r3
addi r3, 1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
sub p5, p2 → p8
```

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

24

24

Renaming example

Original insns

```
xor r1, r2 - r3
add r3, r4 - r4
sub r5, r2 - r3
addi r3, 1 - r1
```

Renamed insns

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 -
```

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

25

25

Renaming example

Original insns

```
xor r1, r2 - r3
add r3, r4 - r4
sub r5, r2 - r3
addi r3, 1 - r1
```

Renamed insns

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

26

26

Renaming example

Original insns

```
xor r1, r2 - r3
add r3, r4 - r4
sub r5, r2 - r3
addi r3, 1 - r1
```

Renamed insns

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

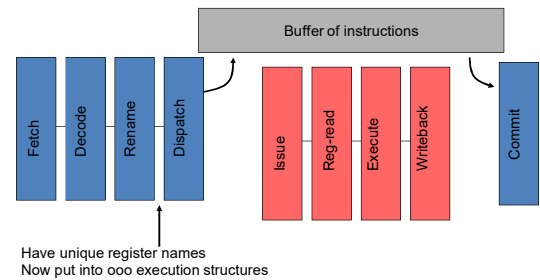
p10

Free-list

27

27

Out-of-order Pipeline



28

28

DYNAMIC SCHEDULING

29

29

Dispatch


- Renamed instructions into ooo structures
 - Re-order buffer (ROB)
 - Holds all instructions until they commit
 - Issue Queue
 - Un-executed instructions
 - Central piece of scheduling logic
 - Content Addressable Memory (CAM) (more later)

30

30

Issue Queue

- Holds un-executed instructions
- Tracks ready inputs
 - Physical register names + ready bit
 - AND to tell if ready



31

31

Dispatch Steps

- Allocate IQ slot
 - Full? Stall
- Read **ready bits** of inputs
 - Table 1-bit per preg
- Clear **ready bit** of output in table
 - Instruction has not produced value yet
- Write data in IQ slot

32

32

Dispatch Example

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Age

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	y
p8	y
p9	y

33

33

Dispatch Example

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	y
p8	y
p9	y

34

34

Dispatch Example

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	y
p9	y

35

35

Dispatch Example

```
xor p1, p2 - p6
add p6, p4 - p7
sub p5, p2 - p8
addi p8, 1 - p9
```

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	y

36

36

Dispatch Example

```
xor p1, p2 -> p6
add p6, p4 -> p7
sub p5, p2 -> p8
addi p8, 1 -> p9
```

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	n

Issue Queue


Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

37

37

Out-of-order pipeline

- Execution (ooo) stages
- Select** ready instructions
 - Send for execution
- Wakeup** dependents



38

38

Dynamic Scheduling/Issue Algorithm

- Data structures:
 - Ready table[phys_reg] → yes/no (part of issue queue)
- Algorithm at "schedule" stage (prior to read registers):


```
foreach instruction:
    if table[insn.phys_input1] == ready &&
       table[insn.phys_input2] == ready then
        insn is "ready"
select the oldest "ready" instruction
table[insn.phys_output] = ready
```

39

39

Issue = Select + Wakeup

- Select** N oldest, ready instructions
 - N=1, "xor"
 - N=2, "xor" and "sub"
 - Note: may have execution resource constraints: *i.e.*, load/store/fp

Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

Ready!

Ready!

40

40

Issue = Select + Wakeup

- Wakeup** dependent instructions
 - CAM search for Dst in inputs
 - Set ready
 - Also update ready-bit table for future instructions

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	n
p8	y
p9	n

Insn	Inp1	R	Inp2	R	Dst	Age
xor	p1	y	p2	y	p6	0
add	p6	y	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	y	---	y	p9	3

41

41

Issue

- Select/Wakeup** one cycle
- Dependents go back to back
 - Next cycle: add/addi are ready:

Insn	Inp1	R	Inp2	R	Dst	Age
add	p6	y	p4	y	p7	1
addi	p8	y	---	y	p9	3

42


42

Register Read

- When do instructions read the register file?
- Option #1: after select, right before execute
 - (Not done at decode)
 - Read **physical** register (renamed)
 - Or get value via bypassing (based on physical register name)
 - This is Pentium 4, MIPS R10k, Alpha 21264 style
- Physical register file may be large
 - Multi-cycle read
- Option #2: as part of issue, keep values in Issue Queue
 - Pentium Pro, Core 2, Core i7


43

OOO execution (2-wide)




48

OOO execution (2-wide)




49

OOO execution (2-wide)




50

OOO execution (2-wide)



51

OOO execution (2-wide)



52

43

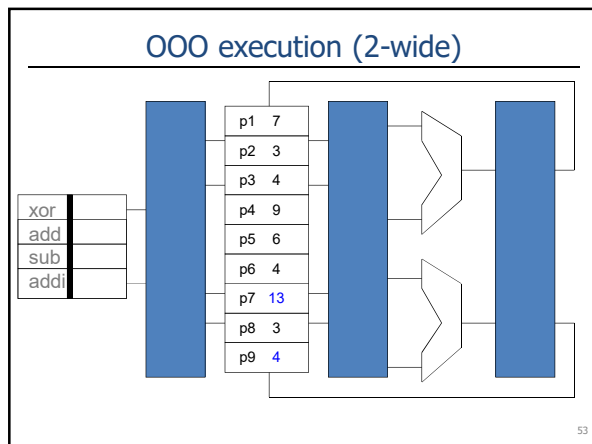
48

49


50

51

52



53



54

Multi-cycle operations

- Multi-cycle ops (load, fp, multiply, *etc.*)
 - Wakeup deferred a few cycles
 - Structural hazard?
- Cache misses?
 - Speculative wake-up (assume hit)
 - Cancel exec of dependents
 - Re-issue later
 - Details: complicated, not important

55

55

Re-order Buffer (ROB)

- All instructions in order
- Two purposes
 - Misprediction recovery
 - In-order commit
 - Maintain appearance of in-order execution
 - Freeing of physical registers

56

56

RENAMING REVISITED

57

57

Renaming revisited

- Overwritten register
 - Freed at commit
- Restore in map table on recovery
 - Branch mis-prediction recovery
- Also must be read at rename

58

58

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

59

59

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

Renamed insns

```
xor p1, p2 →
```

Overwritten Reg

[p3]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

60

60

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

Renamed insns

```
xor p1, p2 → p6
```

Overwritten Reg

[p3]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

61

61

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 →
```

Overwritten Reg

[p3]

[p4]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

62

62

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
```

Overwritten Reg

[p3]

[p4]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

63

63

Renaming example

Original insns

```
xor r1,r2 → r3
add r3,r4 → r4
sub r5,r2 → r3
addi r3,1 → r1
```

Renamed insns

```
xor p1, p2 → p6
add p6, p4 → p7
sub p5, p2 →
```

Overwritten Reg

[p3]

[p4]

[p6]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

64

64

Renaming example

Original insns	Renamed insns	Overwritten Reg
xor r1,r2 → r3	xor p1, p2 → p6	[p3]
add r3,r4 → r4	add p6, p4 → p7	[p4]
sub r5,r2 → r3	sub p5, p2 → p8	[p6]
addi r3,1 → r1		

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

65

65

Renaming example

Original insns	Renamed insns	Overwritten Reg
xor r1,r2 → r3	xor p1, p2 → p6	[p3]
add r3,r4 → r4	add p6, p4 → p7	[p4]
sub r5,r2 → r3	sub p5, p2 → p8	[p6]
addi r3,1 → r1	addi p8, 1 →	[p1]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

66

66

Renaming example

Original insns	Renamed insns	Overwritten Reg
xor r1,r2 → r3	xor p1, p2 → p6	[p3]
add r3,r4 → r4	add p6, p4 → p7	[p4]
sub r5,r2 → r3	sub p5, p2 → p8	[p6]
addi r3,1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10

Free-list

67

67

ROB

- ROB entry holds all info for recover/commit
 - Logical register names
 - Physical register names
 - Instruction types
- Dispatch: insert at tail
 - Full? Stall
- Commit: remove from head
 - Not completed? Stall

68

68

Recovery

- Completely remove wrong path instructions
 - Flush from IQ
 - Remove from ROB
 - Restore map table to before misprediction
 - Free destination registers

69

69

Recovery example

Original insns	Renamed insns	Overwritten Reg
bnz r1, loop	bnz p1, loop	[]
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10

Free-list

70

70

Original insns	Renamed insns	Overwritten Reg
bznz r1, loop	bznz p1, loop	[]
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

71

<u>Original insns</u>	<u>Renamed insns</u>	<u>Overwritten Reg</u>
bnz r1, loop	bnz p1, loop	[]
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

72

<u>Original insns</u>	<u>Renamed insns</u>	<u>Overwritten Reg</u>
bnz r1, loop	bnz p1, loop	[]
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

73

<u>Original insns</u>	<u>Renamed insns</u>	<u>Overwritten Reg</u>
bnz r1, loop	bnz p1, loop	[]
xor r1, r2 → r3	xor p1, p2 → p6	[p3]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

74

<u>Original insns</u>	<u>Renamed insns</u>	<u>Overwritten Reg</u>
b nz r1 loop	b nz p1, loop	[]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

75

- Stores: Write D\$, not registers
 - Can we rename memory?
 - Recover in the cache?

No (at least not easily)

- Cache writes unrecoverable
- Stores: only when certain
 - Commit

- Stores: Write D\$, not registers
 - Can we rename memory?
 - Recover in the cache?
- No (at least not easily)
- Cache writes unrecoverable
 - Stores: only when certain
 - Commit

No (at least not easily)

76

Commit

Original insns	Renamed insns	Overwritten Reg
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

- At commit: instruction becomes architected state
- In-order
- Only when instructions are finished
- Free overwritten register (why?)

77

Freeing over-written register

Original insns	Renamed insns	Overwritten Reg
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

- Before xor: r3 → p3
- After xor: r3 → p6
 - Insns older than xor read p3
 - Insns younger than xor read p6 (until next r3-writing instruction)
- At commit of xor, no older instructions exist
 - No one else needs p3 → free it!

78

Commit Example

Original insns	Renamed insns	Overwritten Reg
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10

Free-list

79

Commit Example

Original insns	Renamed insns	Overwritten Reg
xor r1, r2 → r3	xor p1, p2 → p6	[p3]
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3

Free-list

80

Commit Example

Original insns	Renamed insns	Overwritten Reg
add r3, r4 → r4	add p6, p4 → p7	[p4]
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4

Free-list

81

Commit Example

Original insns	Renamed insns	Overwritten Reg
sub r5, r2 → r3	sub p5, p2 → p8	[p6]
addi r3, 1 → r1	addi p8, 1 → p9	[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6

Free-list

82

77

78

79

80

81

82

Commit Example

Original insns

Renamed insns

Overwritten Reg

addi r3,1 → r1

addi p8, 1 → p9

[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

83

83

Out of order pipeline diagrams


- Standard style: large and cumbersome
- Change layout slightly
 - Columns = stages (dispatch, issue, *etc.*)
 - Rows = instructions
 - Content of boxes = cycles
- For our purposes: issue/exec = 1 cycle
 - Ignore preg read latency, *etc.*
 - Load-use, mul, div, and FP longer

84

84

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2				
add p2, p3 → p4				
xor p4, p5 → p6				
ld [p7] → p8				



85

85

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2				
add p2, p3 → p4				
xor p4, p5 → p6				
ld [p7] → p8				

2-wide
Infinite ROB, IQ, Pregs
Loads: 3 cycles

86

86

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1			
add p2, p3 → p4	1			
xor p4, p5 → p6				
ld [p7] → p8				

Cycle 1:

- Dispatch xor and ld

87

87

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	
add p2, p3 → p4	1			
xor p4, p5 → p6	2			
ld [p7] → p8	2			

Cycle 2:

- Dispatch xor and ld
- 1st Ld issues -- also note WB cycle while you do this (Note: don't issue if WB ports full)

88

88

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	
add p2, p3 → p4	1			
xor p4, p5 → p6	2			
ld [p7] → p8	2	3	6	

Cycle 3:

- add and xor are not ready
- 2nd load is → issue it

89

89

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	
add p2, p3 → p4	1	5	6	
xor p4, p5 → p6	2			
ld [p7] → p8	2	3	6	

Cycle 4:

- nothing

Cycle 5:

- add can issue

90

90

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	6
add p2, p3 → p4	1	5	6	
xor p4, p5 → p6	2	6	7	
ld [p7] → p8	2	3	6	

Cycle 6:

- 1st load can commit (oldest instruction & finished)
- xor can issue

91

91

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	6
add p2, p3 → p4	1	5	6	7
xor p4, p5 → p6	2	6	7	
ld [p7] → p8	2	3	6	

Cycle 7:

- add can commit (oldest instruction & finished)

92

92

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	6
add p2, p3 → p4	1	5	6	7
xor p4, p5 → p6	2	6	7	8
ld [p7] → p8	2	3	6	8

Cycle 8:


- xor and ld can commit (2-wide: can do both at once)

93

93

Out of order pipeline diagrams

Instruction	Disp	Issue	WB	Commit
ld [p1] → p2	1	2	5	6
add p2, p3 → p4	1	5	6	7
xor p4, p5 → p6	2	6	7	8
ld [p7] → p8	2	3	6	8



94

94

HANDLING MEMORY OPS

95

Dynamically Scheduling Memory Ops

- Compilers must schedule memory ops conservatively
- Options for hardware:
 - Hold loads until all prior stores execute (conservative)
 - Execute loads as soon as possible, detect violations (aggressive)
 - When a store executes, it checks if any later loads executed too early (to same address). If so, flush pipeline
 - Learn violations over time, selectively reorder (predictive)

Before Wrong(?)

```
ld r2,4(sp)      ld r2,4(sp)
ld r3,8(sp)      ld r3,8(sp)
add r3,r2,r1      ld r5,0(r8) //does r8==sp?
st r1,0(sp)       add r3,r2,r1
ld r5,0(r8)       ld r6,4(r8) //does r8+4==sp?
ld r6,4(r8)       st r1,0(sp)
sub r5,r6,r4      sub r5,r6,r4
st r4,8(r8)       st r4,8(r8)
```

//stall

96

Loads and Stores

Instruction	Disp	Issue	WB	Commit
fdiv p1,p2 → p3	1	2	25	
st p4 → [p5]	1	2	3	
st p3 → [p6]	2			
ld [p7] → p8	2			

Cycle 3:

- Can ld [p7]→p8 execute? (why or why not?)

97

Loads and Stores

Instruction	Disp	Issue	WB	Commit
fdiv p1,p2 → p3	1	2	25	
st p4 → [p5]	1	2	3	
st p3 → [p6]	2			
ld [p7] → p8	2			

Aliasing (again)

- p5 == p7 ?
- p6 == p7 ?

98

Loads and Stores

Instruction	Disp	Issue	WB	Commit
fdiv p1,p2 → p3	1	2	25	
st p4 → [p5]	1	2	3	
st p3 → [p6]	2			
ld [p7] → p8	2			

Suppose p5 == p7 and p6 != p7

- Can ld [p7]→p8 execute? (why or why not?)

99

Memory Forwarding

- Stores write cache at commit
 - Commit is in-order, delayed by all instructions
 - Allows stores to be "undone" on branch mis-predictions, etc.
- Loads read cache
 - Early execution of loads is critical
- Forwarding
 - Allow store → load communication before store commit
 - Conceptually like reg. bypassing, but different implementation
 - Why? Addresses unknown until execute

100

95

96

97

98

99

100

Forwarding: Store Queue

Store Queue


- Holds all in-flight stores
- CAM: searchable by address
- Age logic: determine youngest matching store older than load

Store execution

- Write Store Queue
 - Address + Data

Load execution

- Search SQ
 - Match? Forward
- Read D\$



101

101

Load scheduling

Store→Load Forwarding:

- Get value from executed (but not committed) store to load
- Load Scheduling:
 - Determine when load can execute with regard to older stores

Conservative load scheduling:

- All older stores have executed
- Some architectures: split store address / store data
 - Only require known address
- Advantage: always safe
- Disadvantage: performance (limits out-of-orderness)

102

102

Our example from before

```
ld [r1] → r5
ld [r2] → r6
add r5, r6 → r7
st r7 → [r3]
ld 4[r1] → r5
ld 4[r2] → r6
add r5, r6 → r7
st r7 → 4[r3]
// loop control here
```

With conservative load scheduling,
what can go out of order?

103

103

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1			
2	ld [p2] → p6	1			
3	add p5, p6 → p7				
4	st p7 → [p3]				
5	ld 4[p1] → p8				
6	ld 4[p2] → p9				
7	add p8, p9 → p4				
8	st p4 → 4[p3]				

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 1:
Dispatch insns #1, #2

104

104

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	
2	ld [p2] → p6	1			
3	add p5, p6 → p7	2			
4	st p7 → [p3]	2			
5	ld 4[p1] → p8				
6	ld 4[p2] → p9				
7	add p8, p9 → p4				
8	st p4 → 4[p3]				

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 2:
Why don't we issue #2?

105

105

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	
2	ld [p2] → p6	1	3	6	
3	add p5, p6 → p7	2			
4	st p7 → [p3]	2			
5	ld 4[p1] → p8	3			
6	ld 4[p2] → p9	3			
7	add p8, p9 → p4				
8	st p4 → 4[p3]				

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 3:
Why don't we issue #3?
Why don't we issue #4?

106

106

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	
2	ld [p2] → p6	1	3	6	
3	add p5,p6 → p7	2			
4	st p7 → [p3]	2			
5	ld 4[p1] → p8	3			
6	ld 4[p2] → p9	3			
7	add p8,p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 4:
Why don't we issue #5?

107

107

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	
3	add p5,p6 → p7	2	6	7	
4	st p7 → [p3]	2			
5	ld 4[p1] → p8	3			
6	ld 4[p2] → p9	3			
7	add p8,p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 6:
Finally some action!

108

108

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5,p6 → p7	2	6	7	
4	st p7 → [p3]	2	7	8	
5	ld 4[p1] → p8	3			
6	ld 4[p2] → p9	3			
7	add p8,p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 7:
Getting somewhere....

109

109

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5,p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	
5	ld 4[p1] → p8	3	8	11	
6	ld 4[p2] → p9	3			
7	add p8,p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 8:
Etc...

110

110

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5,p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	
6	ld 4[p2] → p9	3	9	12	
7	add p8,p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 9:
Etc...

111

111

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5,p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	12
6	ld 4[p2] → p9	3	9	12	
7	add p8,p9 → p4	4	12	13	
8	st p4 → 4[p3]	4			

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 12:
Yawn...

112

112

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5, p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	12
6	ld 4[p2] → p9	3	9	12	13
7	add p8, p9 → p4	4	12	13	
8	st p4 → 4[p3]	4	13	14	

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 13:
Stretch...

113

113

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5, p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	12
6	ld 4[p2] → p9	3	9	12	13
7	add p8, p9 → p4	4	12	13	14
8	st p4 → 4[p3]	4	13	14	

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 14:
Zzzzzz...

114

114

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5, p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	12
6	ld 4[p2] → p9	3	9	12	13
7	add p8, p9 → p4	4	12	13	14
8	st p4 → 4[p3]	4	13	14	15

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 15:
2-wide ooo = 1-wide inorder
I am going to cry.

115

115

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5, p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	8	11	12
6	ld 4[p2] → p9	3	9	12	13
7	add p8, p9 → p4	4	12	13	14
8	st p4 → 4[p3]	4	13	14	15

- 2 wide, conservative scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

What was #5 waiting for?
Can I speculate?

116

116

Load Speculation


- Speculation requires two things....
 - Detection of mis-speculations
 - How can we do this?
 - Recovery from mis-speculations
 - Squash from offending load
 - Saw how to squash from branches: same method

117

117

Load Queue

- Detects ld ordering violations
- Execute load: write addr to LQ
 - Also note any store forwarded from
- Execute store: search LQ
 - Younger load with same addr?
 - Didn't forward from younger store?



118

118

Store Queue + Load Queue

- Store Queue: handles forwarding
 - Written by stores (@ execute)
 - Searched by loads (@ execute)
 - Read SQ when you write to the data cache (@ commit)
- Load Queue: detects ordering violations
 - Written by loads (@ execute)
 - Searched by stores (@ execute)
- Both together
 - Allows aggressive load scheduling
 - Stores don't constrain load execution

119

119

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	
2	ld [p2] → p6	1	3	6	
3	add p5, p6 → p7	2			
4	st p7 → [p3]	2			
5	ld 4[p1] → p8	3	4	7	
6	ld 4[p2] → p9	3			
7	add p8, p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, **aggressive** scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 4:
Speculatively execute #5
before the store (#4).

120

120

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	
2	ld [p2] → p6	1	3	6	
3	add p5, p6 → p7	2			
4	st p7 → [p3]	2			
5	ld 4[p1] → p8	3	4	7	
6	ld 4[p2] → p9	3	5	8	
7	add p8, p9 → p4	4			
8	st p4 → 4[p3]	4			

- 2 wide, **aggressive** scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Cycle 5:
Speculatively execute #6
before the store (#4).

121

121

Our example from before

		Disp	Issue	WB	Commit
1	ld [p1] → p5	1	2	5	6
2	ld [p2] → p6	1	3	6	7
3	add p5, p6 → p7	2	6	7	8
4	st p7 → [p3]	2	7	8	9
5	ld 4[p1] → p8	3	4	7	9
6	ld 4[p2] → p9	3	5	8	10
7	add p8, p9 → p4	4	8	9	10
8	st p4 → 4[p3]	4	9	10	11

- 2 wide, **aggressive** scheduling
- issue 1 load per cycle
- loads take 3 cycles to complete

Fast forward:
4 cycles faster
Actually ooo this time!

122

122

Aggressive Load Scheduling

- Allows loads to issue before older stores
 - Increases out-of-orderness
 - + When no conflict, increases performance
 - Conflict → squash → worse performance than waiting
- Some loads might forward from stores
 - Always aggressive will squash a lot
- Can we have our cake AND eat it too?

123

123

Predictive Load Scheduling

- Predict which loads must wait for stores
- Fool me once, shame on you—fool me twice?
 - Loads default to aggressive
 - Keep table of load PCs that have been caused squashes
 - Schedule these conservatively
 - + Simple predictor
 - Makes “bad” loads wait for *all* older stores: not great
- More complex predictors used in practice
 - Predict which stores loads should wait for

124

124

Out of Order: Window Size

- Scheduling scope = ooo window size
 - Larger = better
- Constrained by physical registers (#preg)
 - ROB roughly limited by #preg = ROB size + #logical registers
 - Big register file = hard/slow
- Constrained by issue queue
 - Limits number of un-executed instructions
 - CAM = can't make big (power + area)
- Constrained by load + store queues
 - Limit number of loads/stores
 - CAMs
 - Active area of research: scaling window sizes
- Usefulness of large window: limited by branch prediction
 - 95% branch mis-prediction rate: 1 in 20 branches, 1 in 100 insns

125

125

Out of Order: Benefits

- Allows speculative re-ordering
 - Loads / stores
 - Branch prediction
- Schedule can change due to cache misses
 - Different schedule optimal from on cache hit
- Done by hardware
 - Compiler may want different schedule for different hw configs
 - Hardware has only its own configuration to deal with

126

126

Static vs. Dynamic Scheduling

- If we can do this in software...
- ...why build complex (slow-clock, high-power) hardware?
 - + Performance portability
 - Don't want to recompile for new machines
 - + More information available
 - Memory addresses, branch directions, cache misses
 - + More registers available
 - Compiler may not have enough to schedule well
 - + Speculative memory operation re-ordering
 - Compiler must be conservative, hardware can speculate
 - But compiler has a larger scope
 - Compiler does as much as it can (not much)
 - Hardware does the rest

127

127

Out of Order: Top 5 Things to Know

- Register renaming
 - How to perform it and how to recover it
- Commit
 - Precise state (ROB)
 - How/when registers are freed
- Issue/Select
 - Wakeup: CAM
 - Choose N oldest ready instructions
- Stores
 - Write at commit
 - Forward to loads via SQ
- Loads
 - Conservative/aggressive/predictive scheduling
 - Violation detection via LQ

128

128

Summary: Dynamic Scheduling

- Dynamic scheduling
 - Totally in the hardware
 - Also called "out-of-order execution" (OoO)
- Fetch many instructions into instruction window
 - Use branch prediction to speculate past (multiple) branches
 - Flush pipeline on branch misprediction
- Rename to avoid false dependencies
- Execute instructions as soon as possible
 - Register dependencies are known
 - Handling memory dependencies more tricky
- "Commit" instructions in order
 - Anything strange happens pre-commit, just flush the pipeline
- Current machines: 100+ instruction scheduling window

129

129