


# CSE 560 Computer Systems Architecture

## Virtual Memory

1

## This Unit: Virtual Memory



- The operating system (OS)
  - A super-application
  - Hardware support for an OS
- Virtual memory
  - Page tables and address translation
  - TLBs and memory hierarchy issues

2

## A Computer System: Hardware


CPUs and memories

- Connected by memory bus

**I/O peripherals:** storage, input, display, network, ...

(NIC = Network Interface Controller)


- With separate or built-in DMA (direct memory access)
- Connected by **system bus** (which is connected to memory bus)



3

## A Computer System: + App Software

- **Application software:** computer must do something




4

## A Computer System: + OS

**Operating System (OS):** virtualizes hardware for apps

- **Abstraction:** provides **services** (e.g., threads, files, etc.)
  - + Simplifies app programming model, raw hardware is nasty
- **Isolation:** gives each app illusion of private CPU, memory, I/O
  - + Simplifies app programming model
  - + Increases hardware resource utilization



5

## Operating System (OS) and User Apps

- Same system development requires a split
- **Operating System (OS):** a super-privileged process
  - Manages hw resource allocation/revocation for all processes
  - Has direct access to resource allocation features
  - **Aware of:** many nasty hardware details, other processes
  - Talks directly to input/output devices (device driver software)
- **User-level apps:** ignorance is bliss
  - **Unaware of:** most nasty hardware details, other apps, OS
  - Explicitly denied access to resource allocation features

6

## System Calls

**System Call:** a user-level app "function call" to OS

- Leave description of what you want done in registers
- SYSCALL instruction (also called TRAP or INT)
  - User-level apps not allowed to invoke arbitrary OS code
  - Restricted set of legal OS addresses to jump to (**trap vector**)

1. Processor jumps to OS via trap vector (begin privileged mode)
2. OS performs operation
3. OS does a "return from system call" (end privileged mode)

7

7

## Interrupts

**Exceptions:** synchronous, generated by running app

- *E.g.*, illegal instruction, divide by zero, *etc.*

**Interrupts:** asynchronous events generated externally

- *E.g.*, timer, I/O request/reply, *etc.*

**Timer:** programmable on-chip interrupt

- Initialize with some number of micro-seconds
- Timer counts down and interrupts when reaches 0

**"Interrupt" handling:** same mechanism for both

- "Interrupts" are on-chip signals/bits
  - Either internal (*e.g.*, timer, exceptions) or from I/O devices
- Processor continuously monitors interrupt status, when true...
- HW jumps to some preset address in OS code (interrupt vector)
- Like an asynchronous, non-programmatic SYSCALL

8

8

## Virtualizing Processors

How do multiple apps (and OS) share the processors?

**Goal:** applications think there are an infinite # of processors

Solution: time-share the resource

- Trigger a **context switch** at a regular interval (~1ms)
  - **Pre-emptive:** app doesn't yield CPU, OS forcibly takes it
    - + Stops greedy apps from starving others
- **Architected state:** PC, registers
  - Save and restore them on context switches
  - Memory state?
- **Non-architected state:** caches, predictor tables, *etc.*
  - Ignore or flush
- Operating System responsible for handling context switching
  - Hardware support is just a timer interrupt

9

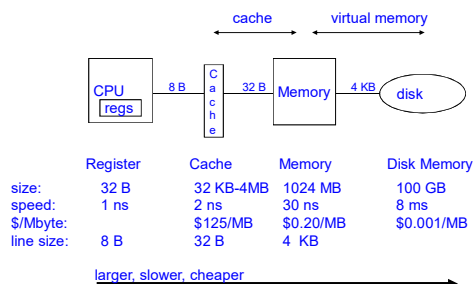
9

## Motivations for Virtual Memory

- **Use Physical DRAM as a Cache for the Disk**
  - Address space of a process can exceed physical memory size
  - Sum of address spaces of multiple processes can exceed physical memory
- **Simplify Memory Management**
  - Multiple processes resident in main memory
    - Each process with its own address space
  - Only "active" code and data is actually in memory
    - Allocate more memory to process as needed
- **Provide Protection**
  - One process can't interfere with another
    - because they operate in different address spaces
  - User process cannot access privileged information
    - different sections of address spaces have different permissions

10

## Levels in Memory Hierarchy



11

## Virtualizing Main Memory

How do multiple apps (and the OS) share main memory?

**Goal:** each application thinks it has private memory

App's insn/data footprint > main memory ?

- **Requires main memory to act like a cache**
  - With disk as next level in memory hierarchy (slow)
  - Write-back, write-allocate, large blocks or "pages"

Solution:

- Part #1: treat memory as a "cache"
- Part #2: add a level of indirection (address translation)

Parameter	I\$/D\$	L2	Main Memory
$t_{hit}$	2ns	10ns	30ns
$t_{miss}$	10ns	30ns	10ms (10M ns)
Capacity	8-64KB	128KB-2MB	64MB-64GB
Block size	16-32B	32-256B	4+KB
Assoc./Repl.	1-4, NMRU	4-16, NMRU	Full, "working set"


12

12

## A System with Physical Memory Only

### Examples:

- most Cray machines, early PCs, many embedded systems, etc.




Addresses generated by the CPU correspond directly to bytes in physical memory

13

## A System with Virtual Memory

### Examples:

- workstations, servers, modern PCs, etc.




Address Translation: Hardware converts virtual addresses to physical addresses via OS-managed lookup table (page table)

14


## Page Faults (like "Cache Misses")

- What if an object is on disk rather than in memory?
  - Page table entry indicates virtual address not in memory
  - OS exception handler invoked to move data from disk into memory
    - current process suspends, others can resume
    - OS has full control over placement, etc.

Before fault




After fault



15

## Virtual Memory (VM)



- Programs use **virtual addresses (VA)**
  - $0 \dots 2^N - 1$
  - VA size also referred to as machine size
  - E.g., 32-bit (embedded) or 64-bit (server)
- Memory uses **physical addresses (PA)**
  - $0 \dots 2^M - 1$  (typically  $M < N$ , especially if  $N = 64$ )
  - $2^M$  is most physical memory machine supports
- VA  $\rightarrow$  PA at **page** granularity (VP  $\rightarrow$  PP)
  - By "system" (OS + HW)
  - Mapping need not preserve contiguity
  - VP need not be mapped to any PP
  - Unmapped VPs live on disk (swap)


16

16

## Virtual Memory (VM)

### Virtual Memory (VM):

- Level of indirection
- Application generated addresses are **virtual addresses (VAs)**
  - Each process **thinks** it has its own  $2^N$  bytes of address space
- Memory accessed using **physical addresses (PAs)**
- VAs translated to PAs at some coarse granularity
- OS controls VA to PA mapping for itself and all other processes
- Logically: translation performed before every insn fetch, load, store
- Physically: hardware acceleration removes translation overhead




17

17

## Uses of Virtual Memory

- Isolation and Multi-programming (Memory Management)**
  - Each app thinks it has  $2^N$  B of memory that starts @ 0
  - Apps can't read/write each other's memory
    - Can't even address the other program's memory!
- Protection**
  - Each page has read/write/execute permission set by OS
  - Enforced by hardware
- Inter-process communication**
  - Map same physical pages into multiple virtual address spaces
  - Or share files via the UNIX `mmap()` call




18

18

## Memory Management


- Multiple processes can reside in physical memory.
- How do we resolve address conflicts?
  - what if two processes access something at the same address?



19

## Solution: Separate Virt. Addr. Spaces

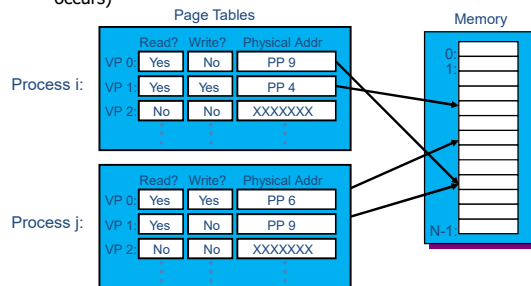
- Virtual and physical address spaces divided into equal-sized blocks
  - blocks are called "pages" (both virtual and physical)
- Each process has its own virtual address space
  - operating system controls how virtual pages as assigned to physical memory



20

## Protection


- Page table entry contains access rights information
  - hardware enforces this protection (trap into OS if violation occurs)



21

## Virtual Memory: The Basics

- Programs use **virtual addresses (VA)**
  - VA size (N) aka machine size (e.g., Core 2 Duo: 48-bit)
- Memory uses **physical addresses (PA)**
  - PA size (M) typically  $M < N$ , especially if  $N=64$
  - $2^M$  is most physical memory machine supports
- VA → PA at **page** granularity (VP → PP)
  - Mapping need not preserve contiguity
  - VP need not be mapped to any PP
  - Unmapped VPs live on disk (swap) or nowhere (if not yet touched)



22

22