

Roadmap Checkpoint

- Thread-level parallelism (TLP)
- Shared memory model
 - Multiplexed uniprocessor
 - Hardware multithreading
 - Multiprocessing
- Synchronization
 - Lock implementation
 - Locking gotchas
- **Cache coherence**
 - Bus-based protocols
 - Directory protocols
- **Memory consistency models**

38

38

Recall: Simplest Multiprocessor

- What if we don't want to share the L1 caches?
 - Bandwidth and latency issue
- Solution: use per-processor ("private") caches
 - Coordinate them with a **Cache Coherence Protocol**

39

39

Shared-Memory Multiprocessors

Conceptual model

- The shared-memory abstraction
- Familiar and feels natural to programmers
- Life would be easy if systems actually looked like this...

40

40

Shared-Memory Multiprocessors

...but systems actually look more like this

- Processors have caches
- Memory may be physically distributed
- Arbitrary interconnect

41

41

Revisiting Our Motivating Example

Processor 0	Processor 1	CPU0	CPU1	Mem
0: addi r1,accts→r3		\$	\$	
1: ld 0(r3),r4				
2: blt r4,r2,done				
3: sub r4,r2→r4				
4: st r4,0(r3)				
	0: addi r1,accts→r3			
	1: ld 0(r3),r4			
	2: blt r4,r2,done			
	3: sub r4,r2→r4			
	4: st r4,0(r3)			

42

42


No-Cache, No-Problem

Processor 0	Processor 1	CPU0	CPU1	Mem
0: addi r1,accts→r3		\$	\$	
1: ld 0(r3),r4				\$500
2: blt r4,r2,done				\$500
3: sub r4,r2→r4				
4: st r4,0(r3)				
	0: addi r1,accts→r3			\$400
	1: ld 0(r3),r4			\$400
	2: blt r4,r2,done			
	3: sub r4,r2→r4			
	4: st r4,0(r3)			\$300

43

43


Cache Incoherence



- Scenario II(a): processors have write-back caches
 - Potentially 3 copies of **accts[241].bal**: memory, p0\$, p1\$
 - Can get incoherent (inconsistent)

44

Write-Through Doesn't Fix It



- Scenario II(b): processors have write-through caches
 - This time only 2 (different) copies of **accts[241].bal**
 - No problem? What if another withdrawal happens on processor 0?

45

44

45

What To Do?


- No caches?
 - Slow
- Make shared data uncachable?
 - Faster, but still too slow
 - Entire **accts** database is technically "shared"
- Flush all other caches on writes to shared data?
 - May as well not have caches
- Hardware cache coherence**
 - Rough goal: all caches have same data at all times
 - + Minimal flushing, maximum caching → best performance

46

46

Bus-based Multiprocessor

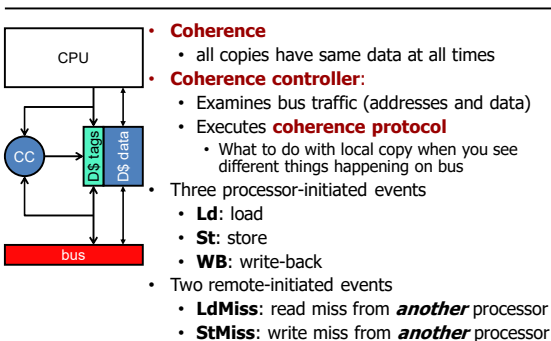
- Simple multiprocessors use a bus
- All processors see **all requests** at the **same time**, same order
- Memory
- Single memory module, **-or-**
 - Banked memory module



47

47

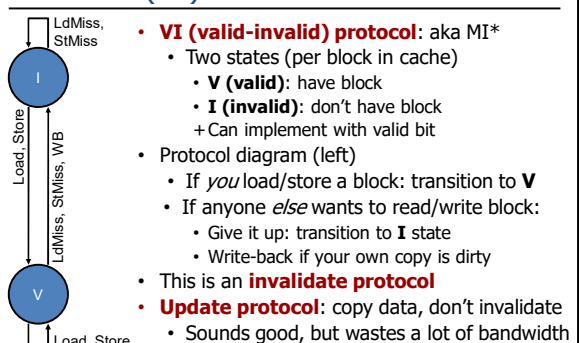
Hardware Cache Coherence



48

48

VI (MI) Coherence Protocol



49

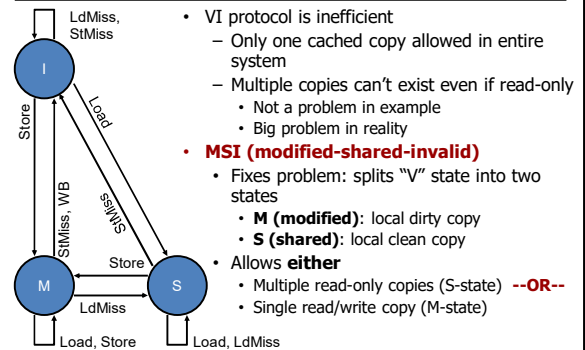
VI Protocol (Write-Back Cache)

Processor 0	Processor 1	CPU0	CPU1	Mem
0: addi r1,accts→r3 1: ld 0(r3),r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)	0: addi r1,accts→r3 1: ld 0(r3),r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)	V:500	V:500	500
		V:400	V:500	500
		I	V:400	400
			V:300	400

- ld by processor 1 generates an "other load miss" event (LdMiss)
- processor 0 responds by sending its dirty copy, transitioning to **I**

50

VI → MSI



51

MSI Protocol (Write-Back Cache)

Processor 0	Processor 1	CPU0	CPU1	Mem
0: addi r1,accts→r3 1: ld 0(r3),r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)	0: addi r1,accts→r3 1: ld 0(r3),r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)	S:500		500
		M:400		500
		S:400	S:400	400
		I	M:300	400

- ld by processor 1 generates a "other load miss" event (LdMiss)
- Processor 0 responds by sending its dirty copy, transitioning to **S**
- st by processor 1 generates a "other store miss" event (StMiss)
- Processor 0 responds by transitioning to **I**


52

Cache Coherence and Cache Misses

- Coherence introduces two new kinds of cache misses
 - Upgrade miss**
 - On stores to read-only blocks
 - Delay to acquire write permission to read-only block
 - Coherence miss**
 - Miss to a block evicted by another processor's requests
- Making the cache larger...
 - Doesn't reduce these type of misses
 - As cache grows large, these sorts of misses dominate
- False sharing**
 - Two or more processors sharing parts of the same block
 - But *not* the same bytes within that block (no actual sharing)
 - Creates pathological "ping-pong" behavior
 - Careful data placement may help, but is difficult

53

MSI → MESI: Exclusive Clean Protocol



- Most modern protocols also include **E (exclusive)** state
- "I have the only cached copy, and it's a **clean** copy"
 - Why is this state useful?

Load transitions to E if no other processors is caching the block, otherwise S

54

Exclusive Clean Protocol Optimization

Processor 0	Processor 1	CPU0	CPU1	Mem
0: addi r1,accts→r3 1: ld 0(r3)→r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)		E:500		500
		M:400		500
		S:400	S:400	400
		I	M:300	400

(No miss)

Processor 1	CPU1	Mem
0: addi r1,accts→r3 1: ld 0(r3),r4 2: blt r4,r2,done 3: sub r4,r2→r4 4: st r4,0(r3)		

55

50

51

52

53

54

55


Snooping Bandwidth Scaling Problems

- Coherence events generated on...
 - L2 misses (and writebacks) – *actually* last level cache misses
- Problem#1: **N^2 bus traffic**
 - All N processors send their misses to all N-1 other processors
 - Assume: 2 IPC, 2 GHz clock, 0.01 misses/insn **per processor**
 - 0.01 misses/insn x 2 insns/cycle x 2 cycle/ns x 64 B blocks = 2.56 GB/s... per processor
 - With 16 processors, that's 40 GB/s! With 128 that's 320 GB/s!!
 - You can use multiple buses... but that hinders global ordering
- Problem#2: **N^2 processor snooping bandwidth**
 - 0.01 events/insn x 2 insns/cycle = 0.02 events/cycle per processor
 - 16 processors: 0.32 bus-side tag lookups per cycle
 - Add 1 extra port to cache tags? Okay
 - 128 processors: 2.56 tag lookups per cycle! 3 extra tag ports?

56

56

"Scalable" Cache Coherence



Part I: **bus bandwidth**

Replace non-scalable bandwidth substrate (bus)...
...with scalable one (point-to-point network, *e.g.*, mesh)

Part II: **processor snooping bandwidth**

- Most snoops result in no action
- Replace non-scalable broadcast protocol (spam everyone)...
...with scalable **directory protocol** (only notify processors that care)

57

57